
[Re] HyperFast: Instant Classification for Tabular Data

Mohamed Amine Kina

Universität Bremen

Msc. Artificial Intelligence and Intelligent Intelligence

mkina@uni-bremen.de

1 Reproducibility Summary

This report documents a replication probe of the “HyperFast” model, specifically focusing on its claims of “instant” inference and competitive performance on small tabular datasets. Using the official pre-trained model artifact, I attempted to replicate the results reported in the paper’s “Mini-test” benchmark (datasets with ≤ 1000 samples).

I successfully verified the inference speed claim: HyperFast generated predictions in 1.5–6.0 seconds per dataset, significantly faster than the 12–113 seconds required to tune a Gradient Boosting baseline. I also validated the performance claim on complex, noisy data: on the Diabetes dataset, HyperFast (73.6% accuracy) outperformed a tuned LightGBM model (72.7%).

However, I encountered a significant “reproducibility gap” on simpler datasets. On Banknote Authentication and Pendigits, HyperFast failed to reproduce the near-perfect accuracy reported in the paper (e.g., scoring 85.2% vs. the reported 100%), performing worse than a standard Logistic Regression baseline. This suggests the model may be brittle or highly sensitive to the specific random downsampling splits used in the original study.

2 Methodology

To conduct a resource-efficient probe, we selected a diverse subset of four datasets from the OpenML-CC18 “Mini-test” suite. My selection covers varied task types and feature modalities: *Diabetes* and *Banknote-authentication* (numerical binary classification), *Pendigits* (numerical multiclass), and *Credit-approval* (mixed-type binary classification).

2.1 Experimental Protocol

Data Processing: I strictly adhered to the paper’s “Mini-test” protocol by performing a stratified 80/20 train–test split and then downsampling the training set to a maximum of 1,000 samples and 100 features. This downsampling is critical to testing the “small data” capabilities claimed by the authors.

Preprocessing: As required for the HyperFast neural network, I applied mean imputation for numerical features, mode imputation for categorical features, and standard scaling/one-hot encoding.

2.2 Baselines

Logistic Regression: Run with default `scikit-learn` settings to establish a performance floor.

LightGBM: Used as the state-of-the-art gradient boosting baseline. To ensure a fair comparison, I utilized `RandomizedSearchCV` with a fixed iteration budget designed to complete within 5 minutes per dataset, matching the time budget allocated in the paper.

Metrics: I reported *Balanced Accuracy* to verify the paper’s findings, which explicitly warn against using standard accuracy for these datasets.

3 Results

The table below compares my replication results against the values reported in Table 7 of the original paper.

Table 1: Comparison of Accuracy: Original Paper vs. Replication Results

Dataset	Log. Reg.		LightGBM		HyperFast	
	Paper	Rep.	Paper	Rep.	Paper	Rep.
Diabetes	66.93 ± 0.0	66.93	71.28 ± 0.33	72.70	70.91 ± 0.0	73.63
Pendigits	92.97 ± 0.59	91.97	96.81 ± 0.48	96.28	98.50 ± 0.38	88.03
Credit Approval	84.35 ± 0.0	83.53	87.00 ± 0.08	82.88	86.44 ± 0.54	81.55
Banknote Auth.	98.69 ± 0.0	97.39	99.84 ± 0.17	99.67	100.0 ± 0.0	85.25

3.1 Analysis

- **Validation of Generalization Potential:** On the *Diabetes* dataset, my probe supports the paper’s central claim regarding zero-shot generalization. HyperFast achieved a higher balanced accuracy (73.6%) than the tuned LightGBM baseline (72.7%) without requiring any gradient updates. This demonstrates that the meta-learning prior can, in specific instances, effectively model unseen real-world tabular data better than traditional boosting methods.
- **Inference Efficiency:** Although I did not strictly achieve sub-second inference due to platform-specific loading overhead (1.5 s–6.2 s), HyperFast remained orders of magnitude faster than the iterative tuning required for LightGBM. For *Pendigits*, HyperFast inference concluded in ~ 1.7 s, whereas the LightGBM hyperparameter search required over 113 s to converge to a comparable solution.
- **Discrepancies on Linearly Separable Tasks:** I observed significant performance degradation on datasets where simple baselines typically excel. On *Banknote Authentication* and *Pendigits*, my replication lagged behind the original paper’s reported scores by –14.8% and –10.5%, respectively. Crucially, on these tasks, even the baseline Logistic Regression outperformed HyperFast in my experiments. This divergence suggests potential sensitivity to feature preprocessing or distribution shifts that the meta-model did not robustly handle in my environment.

4 What Was Easy

- **Model Availability:** The authors provided a pre-trained model artifact (`hyperfast.ckpt`) and a pip-installable package. This allowed me to treat the complex meta-learning framework as a “black box” standard `scikit-learn` classifier (`HyperFastClassifier`), drastically reducing the engineering barrier to entry.
- **Inference Pipeline:** Once the environment was set up, running the model was straightforward. The API consistency with `scikit-learn` meant I could plug HyperFast directly into my existing evaluation loops without writing custom inference code.

5 What Was Difficult

- **The Reproducibility Gap:** The most challenging aspect was diagnosing the root cause of performance degradation on “trivial” datasets like *Banknote Authentication*. Despite adhering to the strict downsampling protocol ($N < 1,000$), my results deviated significantly from the paper’s reported near-perfect scores. This discrepancy highlights a lack of documentation regarding the specific random seeds or stratified splits used in the original study. In small-data regimes, the specific selection of training examples can drastically shift the decision boundary, yet the exact indices were not provided.

- **Resource-Constrained Baseline Comparison:** Reproducing the baselines required a strategic trade-off between fairness and feasibility. While the paper utilizes extensive hyperparameter search spaces, implementing full Bayesian optimization for complex baselines (such as SAINT or TabPFN) was computationally prohibitive within a single-day workflow. Consequently, we had to rely on LightGBM as a representative “efficiency” rival, implementing a time-constrained search to approximate the paper’s experimental conditions.
- **Implicit Preprocessing Requirements:** Although the documentation mentions data imputation, the strict requirement for handling missing values in the neural network backend was understated. Unlike robust tree-based models (e.g., XGBoost), which handle NaN values natively, the HyperFast pipeline is sensitive to input integrity. Ensuring the preprocessing pipeline was perfectly aligned for HyperFast—without leaking information or unfairly penalizing the baselines—required trial-and-error that went beyond the provided “Quickstart” guide.